

1. Execution environment and how to build the programs

Execution environment – Ubuntu 20.04.2 LTS, gcc 9.3.0

```
user@user-VirtualBox:~/53proj$ lsb_release -a
No LSB modules are available.
Distributor ID: Ubuntu
Description:    Ubuntu 20.04.2 LTS
Release:        20.04
Codename:       focal
```

```
user@user-VirtualBox:~/53proj$ gcc --version
gcc (Ubuntu 9.3.0-17ubuntu1~20.04) 9.3.0
Copyright (C) 2019 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

How to build:

server.c – gcc -o server server.c

client.c – gcc -o client client.c

or

make

How to execute

server - ./server <port number>

client - ./client <ip address> <port number>

```
user@user-VirtualBox:~/53proj/proj2$ make
gcc -o client client.c
gcc -o server server.c
user@user-VirtualBox:~/53proj/proj2$ ./server 2000
```

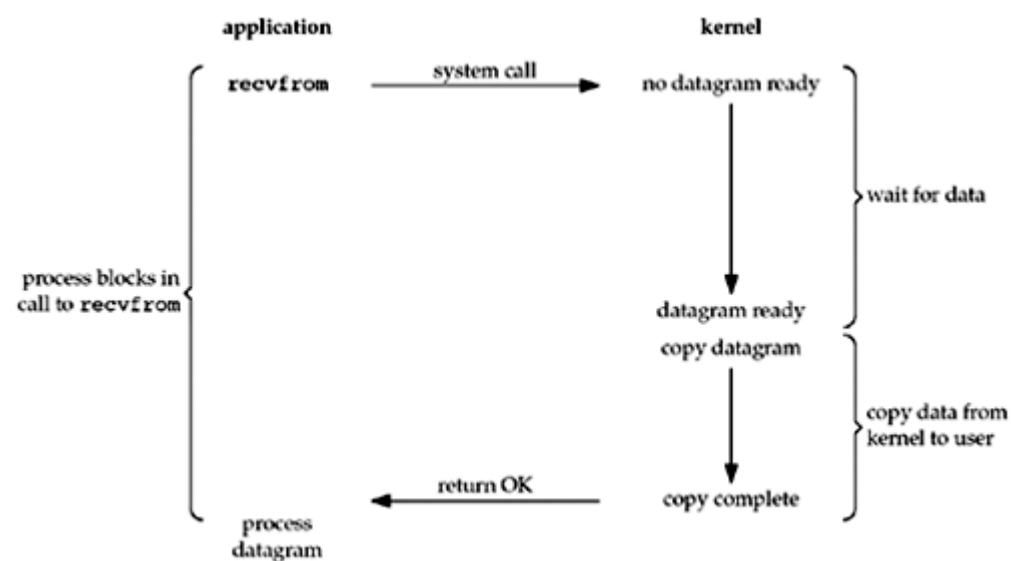
```
user@user-VirtualBox:~/53proj/proj2$ make
gcc -o client client.c
gcc -o server server.c
user@user-VirtualBox:~/53proj/proj2$ ./client 127.0.0.1 2000
```

2. Comparison of the I/O models

-blocking I/O, non-blocking I/O, I/O multiplexing

1) Blocking I/O Model

The most prevalent model for I/O is the *blocking I/O model*, which we have used for all our examples so far in the text. By default, all sockets are blocking. Using a datagram socket for our examples, we have the scenario shown in Figure 1.



<Figure 1> blocking I/O model

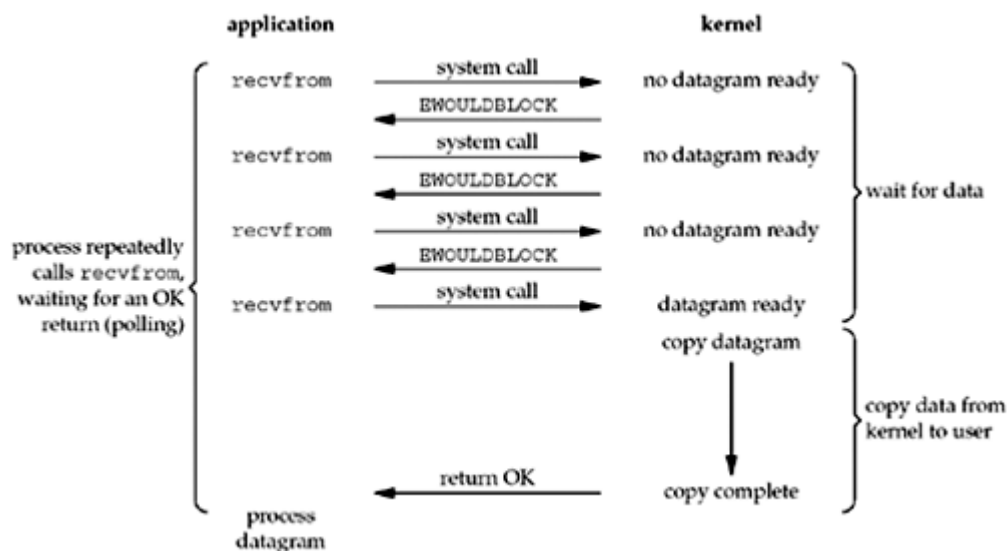
We use UDP for this example instead of TCP because with UDP, the concept of data being "ready" to read is simple: either an entire datagram has been received or it has not. With TCP it gets more complicated, as additional variables such as the socket's low-water mark come into play.

In the examples in this section, we also refer to `recvfrom` as a system call because we are differentiating between our application and the kernel. Regardless of how `recvfrom` is implemented (as a system call on a Berkeley-derived kernel or as a function that invokes the `getmsg` system call on a System V kernel), there is normally a switch from running in the application to running in the kernel, followed at some time later by a return to the application.

In Figure 1, the process calls `recvfrom` and the system call does not return until the datagram arrives and is copied into our application buffer, or an error occurs. The most common error is the system call being interrupted by a signal. We say that our process is *blocked* the entire time from when it calls `recvfrom` until it returns. When `recvfrom` returns successfully, our application processes the datagram.

2) Nonblocking I/O Model

When we set a socket to be nonblocking, we are telling the kernel "when an I/O operation that I request cannot be completed without putting the process to sleep, do not put the process to sleep, but return an error instead.". Figure 2 shows a summary of the example we are considering.



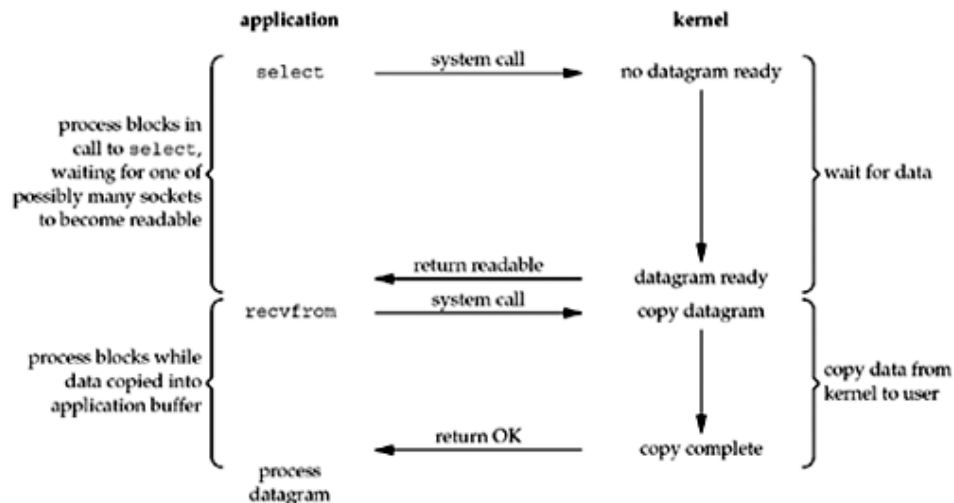
<Figure 2> non-blocking I/O model

The first three times that we call `recvfrom`, there is no data to return, so the kernel immediately returns an error of `EWOULDBLOCK` instead. The fourth time we call `recvfrom`, a datagram is ready, it is copied into our application buffer, and `recvfrom` returns successfully. We then process the data.

When an application sits in a loop calling `recvfrom` on a nonblocking descriptor like this, it is called *polling*. The application is continually polling the kernel to see if some operation is ready. This is often a waste of CPU time, but this model is occasionally encountered, normally on systems dedicated to one function.

3) I/O Multiplexing Model

With *I/O multiplexing*, we call `select` or `poll` and block in one of these two system calls, instead of blocking in the actual I/O system call. Figure 3 is a summary of the I/O multiplexing model.




<Figure 3> I/O multiplexing model

We block in a call to select, waiting for the datagram socket to be readable. When select returns that the socket is readable, we then call recvfrom to copy the datagram into our application buffer.

Comparing Figure 3 to Figure 1, there does not appear to be any advantage, and in fact, there is a slight disadvantage because using select requires two system calls instead of one. But the advantage in using select, which we will see later in this chapter, is that we can wait for more than one descriptor to be ready.

Another closely related I/O model is to use multithreading with blocking I/O. That model very closely resembles the model described above, except that instead of using select to block on multiple file descriptors, the program uses multiple threads (one per file descriptor), and each thread is then free to call blocking system calls like recvfrom.



<Figure 4> comparison of blocking model, nonblocking model, multiplexing I/O model.

3. Comments for each line of the essential part of my code

1)server.c

```
#include <stdio.h>
#include <unistd.h>
#include <sys/socket.h>
#include <sys/types.h>
#include <arpa/inet.h>
#include <stdlib.h>
#include <string.h>
#include <sys/select.h>
#define MAX 1024 //max client number(=FD_SETSIZE)
#define BUF_LEN 1024 //max buffer length
int main(int argc, char *argv[]) {
    //argv error handling
    if(argc!=2){
        printf("usage : %s <port> \n" , argv[0]);
        exit(1);
    }
    //create IPv4, TCP socket
    int sockfd = socket(AF_INET,SOCK_STREAM,0);
    if(sockfd == -1) { //socket create error handler
        printf("create socket fail\n");
        return -1;
    }
    else {
        printf("creat socket success , sockfd = %d\n",sockfd);
    }
    //initialize server address information
    struct sockaddr_in seraddr,cliaddr;
    seraddr.sin_family = AF_INET;
    seraddr.sin_port = htons(atoi(argv[1]));
    seraddr.sin_addr.s_addr = htonl(INADDR_ANY);
    memset(seraddr.sin_zero,0,8);

    //allocate server address
    socklen_t len = sizeof(struct sockaddr);
    int bindret = bind(sockfd,(struct sockaddr *)&seraddr,len);
    if(bindret == -1) { //bind error handler
        printf("bind fail\n");
        close(sockfd);
        return -2;
    }
    else {
        printf("bind success\n");
    }

    //listen for client(max number 1024-FD_SETSIZE)
    int listenret = listen(sockfd,MAX);
    if(listenret == -1) { //listen error handler
        printf("listen fail\n");
        close(sockfd);
        return -3;
    }
    else {
        printf("listen success\n");
    }
    //fd_set, current max file descriptor
    fd_set read,readset;
    int maxfd = sockfd;
    //initialize readset&set sockfd to 1
    FD_ZERO(&readset);
    FD_SET(sockfd,&readset);
```

```

while(1) {
    //clientfd
    int confd;
    read = readset;
    //select which file descriptor has changed
    int selectret = select(maxfd+1,&read,NULL,NULL,NULL);
    if(selectret == -1) { //select error handler
        printf("select fail\n");
        close(sockfd);
        return -4;
    }

    //if socket for tcp request has changed
    if(FD_ISSET(sockfd,&read)) {
        //accept client&allocate confd to client
        confd = accept(sockfd,(struct sockaddr *)&cliaddr,&len);
        printf("connect client(%s) success\nconfd
= %d\n",inet_ntoa(cliaddr.sin_addr),sockfd);
        //update maxfd
        if(maxfd < confd) {
            maxfd = confd;
        }
        //set confd in readset
        FD_SET(confd,&readset);

        //if sockfd is the only file descriptor changed
        if(selectret == 1) {
            continue;
        }
    }
    //message from client
    char buf[BUF_LEN];
    int i;

    //for the rest of file descriptors(client's file descriptor)
    for(i = sockfd+1;i<=maxfd;i++) {
        memset(buf,0,BUF_LEN);

        //if client's file descriptor has changed
        if(FD_ISSET(i,&read)) {
            if(recv(i,buf,BUF_LEN,0) == 0) { //if end of data
                FD_CLR(i,&readset); //clear file descriptor
                printf("-----client(%s) quit-----
\n",inet_ntoa(cliaddr.sin_addr));
                continue;
            }
            else {
                printf("client(%d) send message\n",i);
                send(i,buf,BUF_LEN,MSG_CONFIRM); //echo buf
                if(strncmp(buf,"end",3) == 0) { //end of
                    FD_CLR(i,&readset); //clear file
                    printf("-----client(%s) quit--
-----\n",inet_ntoa(cliaddr.sin_addr));
                    continue;
                }
            }
        }
    }
}
return 0;
}

```

2)client.c

```
#include <stdio.h>
#include <unistd.h>
#include <sys/socket.h>
#include <sys/types.h>
#include <arpa/inet.h>
#include <string.h>
#include <stdlib.h>
#include <sys/select.h>
#define BUF_LEN 1024 // max buffer length
int main(int argc, char *argv[])
{
    //argv error handling
    if (argc != 3) {
        printf("usage : %s <ip> <port> \n" , argv[0]);
        exit(1);
    }

    //create IPv4, TCP socket
    int sockfd = socket(AF_INET,SOCK_STREAM,0);
    if(sockfd < 0) { // socket creation error handler
        printf("creat socket fail\n");
        return -1;
    }
    else {
        printf("creat socket success , sockfd = %d\n",sockfd);
    }

    //initialize server address information
    struct sockaddr_in cliaddr;
    cliaddr.sin_family = AF_INET;
    cliaddr.sin_port = htons(atoi(argv[2]));
    cliaddr.sin_addr.s_addr = inet_addr(argv[1]);
    memset(cliaddr.sin_zero,0,8);
    //connect request
    int confd = connect(sockfd,(struct sockaddr *)&cliaddr,sizeof(struct sockaddr));
    if(confd < 0) { //connection error handler
        printf("connect fail\n");
        close(sockfd);
        return -2;
    }
    else {
        printf("-----connect success-----\nsockfd = %d\n",sockfd);
    }

    // message buffer
    char buf[BUF_LEN];
    int str_len,input_len;

    //fd_set&initialize readset&add stdin/socketfd in readset
    fd_set reads, readset;
    FD_ZERO(&readset);
    FD_SET(sockfd,&readset);
    FD_SET(0,&readset);

    while(1) {
        reads=readset;

        //select which file descriptor has changed
        if(select(sockfd+1,&reads,NULL,NULL,NULL)==-1) { //select error handler
            printf("select() : error\n");
            exit(-1);
        }

        //if stdin has changed
```

```

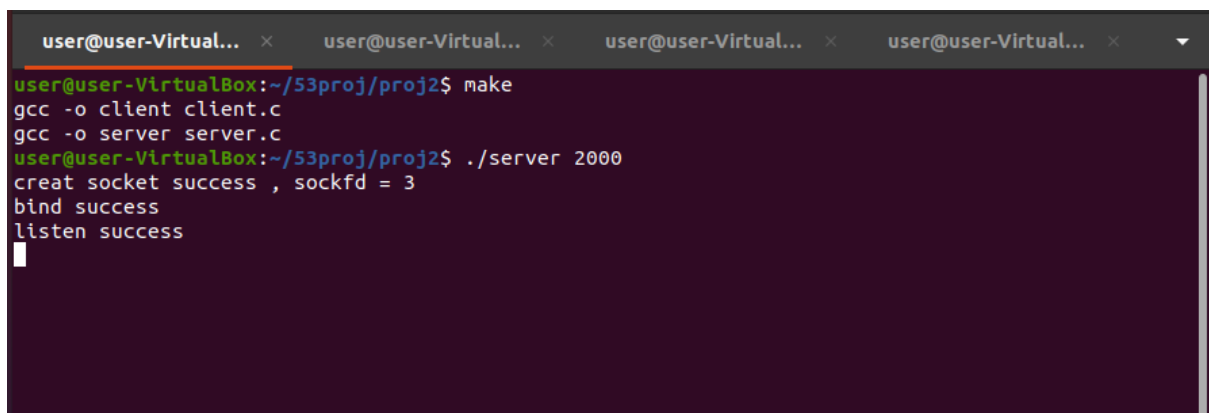
        if(FD_ISSET(0,&reads)) {
            memset(buf,0,BUF_LEN);
            input_len=read(0,buf,BUF_LEN); //read buffer from stdin
            send(sockfd,buf,strlen(buf),0); //send it to the server
        }

        //if sockfd has changed(recv from server)
        if(FD_ISSET(sockfd,&reads)) {
            if(str_len=recv(sockfd,buf,BUF_LEN,0)) { //received buffer from
server
                buf[str_len]='\0';
                printf("echo - ");
                fputs(buf, stdout); // echo it
            }
            else { // no response from server
                FD_CLR(sockfd,&readset); //clear sockfd from readset
                printf("server closed\n");
                close(sockfd);
                break;
            }
        }
    }
    close(sockfd);
    return 0;
}

```

4. Execution result analysis

portnum=2000

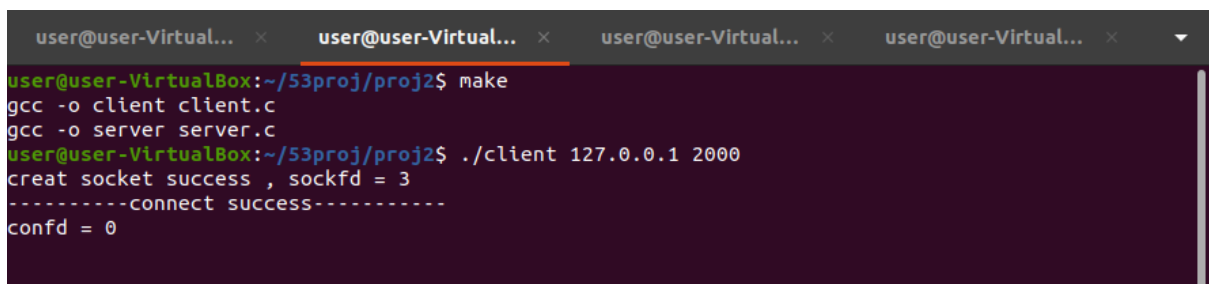


```

user@user-Virtual... x user@user-Virtual... x user@user-Virtual... x user@user-Virtual... x
user@user-VirtualBox:~/53proj/proj2$ make
gcc -o client client.c
gcc -o server server.c
user@user-VirtualBox:~/53proj/proj2$ ./server 2000
creat socket success , sockfd = 3
bind success
listen success

```

<Figure 5> Server – Terminal 1



```

user@user-Virtual... x user@user-Virtual... x user@user-Virtual... x user@user-Virtual... x
user@user-VirtualBox:~/53proj/proj2$ make
gcc -o client client.c
gcc -o server server.c
user@user-VirtualBox:~/53proj/proj2$ ./client 127.0.0.1 2000
creat socket success , sockfd = 3
-----connect success-----
confd = 0

```

<Figure 6> First client – Terminal 2


```
user@user-Virtual... x user@user-Virtual... x user@user-Virtual... x user@user-Virtual... x
user@user-VirtualBox:~/53proj/proj2$ make
gcc -o client client.c
gcc -o server server.c
user@user-VirtualBox:~/53proj/proj2$ ./server 2000
creat socket success , sockfd = 3
bind success
listen success
connect client(127.0.0.1) success
confd = 4
```

<Figure 7> After connection with client1 – Terminal 1

```
user@user-Virtual... x user@user-Virtual... x user@user-Virtual... x user@user-Virtual... x
user@user-VirtualBox:~/53proj/proj2$ ./client 127.0.0.1 2000
creat socket success , sockfd = 3
-----connect success-----
confd = 0
```

<Figure 8> Second client– Terminal 3

```
user@user-Virtual... x user@user-Virtual... x user@user-Virtual... x user@user-Virtual... x
user@user-VirtualBox:~/53proj/proj2$ make
gcc -o client client.c
gcc -o server server.c
user@user-VirtualBox:~/53proj/proj2$ ./server 2000
creat socket success , sockfd = 3
bind success
listen success
connect client(127.0.0.1) success
confd = 4
connect client(127.0.0.1) success
confd = 5
█
```

<Figure 9> Server program after connection of second client – Terminal 1

```
user@user-Virtual... x user@user-Virtual... x user@user-Virtual... x user@user-Virtual... x
user@user-VirtualBox:~/53proj/proj2$ ./client 127.0.0.1 2000
creat socket success , sockfd = 3
-----connect success-----
confd = 0
█
```

<Figure 10> Third client– Terminal 4

```
user@user-VirtualBox: ~/53proj/proj2$ make
gcc -o client client.c
gcc -o server server.c
user@user-VirtualBox:~/53proj/proj2$ ./server 2000
creat socket success , sockfd = 3
bind success
listen success
connect client(127.0.0.1) success
confd = 4
connect client(127.0.0.1) success
confd = 5
connect client(127.0.0.1) success
confd = 6
```

<Figure 9> Server program after connection of third client – Terminal 1

<pre>user@user-VirtualBox:~/53proj/proj2\$./client 127.0.0.1 5000 creat socket success , sockfd = 3 -----connect success----- confd = 0 hheelllloo,, mmyy nnaamee liss yyoouunnggwwoooo kkitnnhh echo - hheelllloo,, mmyy nnaamee liss yyoouunnggwwoooo kkitnnhh echo -</pre>	<pre>user@user-VirtualBox:~/53proj/proj2\$./client 127.0.0.1 5000 creat socket success , sockfd = 3 -----connect success----- confd = 0 hheelllloo,, mmyy nnaamee liss yyoouunnggwwoooo kkitnnhh echo - hheelllloo,, mmyy nnaamee liss yyoouunnggwwoooo kkitnnhh echo -</pre>
<pre>user@user-VirtualBox:~/53proj/proj2\$./client 127.0.0.1 5000 creat socket success , sockfd = 3 -----connect success----- confd = 0 hheelllloo,, mmyy nnaamee liss yyoouunnggwwoooo kkitnnhh echo - hheelllloo,, mmyy nnaamee liss yyoouunnggwwoooo kkitnnhh echo -</pre>	<pre>user@user-VirtualBox:~/53proj/proj2\$./client 127.0.0.1 5000 creat socket success , sockfd = 3 -----connect success----- confd = 0 hello, my name is youngwoo kinh echo - hello, my name is youngwoo kinh</pre>

<Figure 10> Sending message simultaneously using terminator

```
user@user-VirtualBox:~/53proj/proj2$ ./server 5000
creat socket success , sockfd = 3
bind success
listen success
connect client(127.0.0.1) success
confd = 4
connect client(127.0.0.1) success
confd = 5
connect client(127.0.0.1) success
confd = 6
connect client(127.0.0.1) success
confd = 7
client(4) send message
client(6) send message
client(5) send message
client(4) send message
client(6) send message
client(5) send message
client(7) send message
```

<Figure 11> Server program after sending message – Terminal 1

As you can see, the server successfully received simultaneous message from multiple clients. Also, the server and the clients remain online after sending message.